

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call



```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all fi

import os
for dirname, _, filenames in os.walk('/content/drive/MyDrive/Cars Dataset'):
    print(os.path.join(dirname))
```

```
/content/drive/MyDrive/Cars Dataset
/content/drive/MyDrive/Cars Dataset/train
/content/drive/MyDrive/Cars Dataset/train/Hyundai Creta
/content/drive/MyDrive/Cars Dataset/train/Mahindra Scorpio
/content/drive/MyDrive/Cars Dataset/train/Audi
/content/drive/MyDrive/Cars Dataset/train/Rolls Royce
/content/drive/MyDrive/Cars Dataset/train/Toyota Innova
/content/drive/MyDrive/Cars Dataset/train/Swift
/content/drive/MyDrive/Cars Dataset/train/Tata Safari
/content/drive/MyDrive/Cars Dataset/test
/content/drive/MyDrive/Cars Dataset/test/Rolls Royce
/content/drive/MyDrive/Cars Dataset/test/Hyundai Creta
/content/drive/MyDrive/Cars Dataset/test/Swift
/content/drive/MyDrive/Cars Dataset/test/Toyota Innova
/content/drive/MyDrive/Cars Dataset/test/Mahindra Scorpio
/content/drive/MyDrive/Cars Dataset/test/Tata Safari
/content/drive/MyDrive/Cars Dataset/test/Audi
```

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense,Dropout,Convolution2D,MaxPooling2D
import tensorflow as tf
import matplotlib.pyplot as plt
from IPython.display import HTML

```

```

from tensorflow.keras.preprocessing.image import ImageDataGenerator
IMAGE_SIZE = 128

```

```

train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=10,
    horizontal_flip=True
)
train_generator = train_datagen.flow_from_directory(
    '/content/drive/MyDrive/Cars Dataset/train',
    target_size=(IMAGE_SIZE,IMAGE_SIZE),
    class_mode="sparse",
)

```

Found 3352 images belonging to 7 classes.

```

count=0
for image_batch, label_batch in train_generator:
    # print(label_batch)
    print(image_batch[0])
    break

```

```

[[[0.44832808 0.48371983 0.41607684]
  [0.4622689  0.48995802 0.45751396]
  [0.36443672 0.3918877  0.36051515]
  ...
  [0.19959636 0.1914969  0.21619894]
  [0.20328042 0.19354358 0.21947366]
  [0.20696445 0.19559029 0.22274837]]

[[[0.4536495  0.48945057 0.41771418]
  [0.46063155 0.48913932 0.4530112 ]
  [0.37507954 0.40253052 0.37115797]
  ...

```

```
[0.16606945 0.12446638 0.18175572]
[0.16033868 0.11464224 0.17602497]
[0.15460795 0.10481811 0.17029423]]
```

```
[[0.45897087 0.49518132 0.41935155]
 [0.45899418 0.48832065 0.4485085 ]
 [0.38572237 0.41317335 0.3818008 ]
```

...

```
[0.12924187 0.06641182 0.13900332]
[0.12596716 0.06149976 0.13490993]
[0.12269245 0.05658768 0.13081653]]
```

...

```
[[0.8588236 0.8745099 0.87843144]
 [0.8588236 0.8745099 0.87843144]
 [0.8588236 0.8745099 0.87843144]
```

...

```
[0.28235295 0.28627452 0.29411766]
[0.28235295 0.28627452 0.29411766]
[0.28326926 0.28627452 0.29411766]]
```

```
[[0.8588236 0.8745099 0.87843144]
 [0.8588236 0.8745099 0.87843144]
 [0.8588236 0.8745099 0.87843144]
```

...

```
[0.28235295 0.28627452 0.29411766]
[0.28235295 0.28627452 0.29411766]
[0.28285995 0.28627452 0.29411766]]
```

```
[[0.8588236 0.8745099 0.87843144]
 [0.8588236 0.8745099 0.87843144]
 [0.8588236 0.8745099 0.87843144]
```

...

```
[0.28235295 0.28627452 0.29411766]
[0.28235295 0.28627452 0.29411766]
[0.2824506 0.28627452 0.29411766]]]
```

```
class_names = list(train_generator.class_indices.keys())
class_names
```

```
['Audi',
 'Hyundai Creta',
 'Mahindra Scorpio',
```

```
'Rolls Royce',  
'Swift',  
'Tata Safari',  
'Toyota Innova']
```

```
test_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=10,  
    horizontal_flip=True)  
  
test_generator = test_datagen.flow_from_directory(  
    '/content/drive/MyDrive/Cars Dataset/train',  
    target_size=(IMAGE_SIZE,IMAGE_SIZE),  
    class_mode="sparse"  
)
```

Found 3352 images belonging to 7 classes.

```
for image_batch, label_batch in test_generator:  
    print(image_batch[0])  
    break
```

```
[[[0.5441202  0.50553775 0.39838892]  
  [0.7860965  0.7501624  0.6377169 ]  
  [0.868689   0.8333949  0.70657015]  
  ...  
  [0.41640702 0.36998835 0.11791416]  
  [0.4129527  0.3672249  0.11538099]  
  [0.40949836 0.36446142 0.11284781]]  
  
[[[0.55517405 0.51682186 0.4092125 ]  
  [0.7669826  0.7308182  0.61883324]  
  [0.8666164  0.8313223  0.7054187 ]  
  ...  
  [0.46206206 0.41610378 0.17640741]  
  [0.4671284  0.42093986 0.18216461]  
  [0.47219473 0.42577592 0.18792182]]  
  
[[[0.5662279  0.528106  0.42003605]  
  [0.74786866 0.711474  0.5999496 ]  
  [0.86454386 0.82924974 0.7042673 ]  
  ...
```

[0.48064607 0.43358725 0.20392159]
[0.4801855 0.4331267 0.20392159]
[0.4797249 0.4326661 0.20392159]]

...

[[0.5953907 0.57186127 0.52480245]
[0.5965422 0.57301277 0.5259539]
[0.59769356 0.57416415 0.52710533]

...

[0.48627454 0.4666667 0.3921569]
[0.48517397 0.46556613 0.3910563]
[0.48235297 0.46274513 0.38823533]]

[[0.6085648 0.5850353 0.5379765]
[0.60741335 0.58388394 0.5368251]
[0.6062619 0.5827325 0.5356736]

...

[0.48627454 0.4666667 0.3921569]
[0.48540425 0.4657964 0.39128658]
[0.48235297 0.46274513 0.38823533]]

[[0.595323 0.5717936 0.5247348]
[0.59647447 0.57294506 0.52588624]
[0.5976259 0.5740965 0.5270376]

...

[0.48627454 0.4666667 0.3921569]
[0.48563454 0.4660267 0.3915169]
[0.48235297 0.46274513 0.38823533]]]

```
sz = 128
```

```
# Initializing the CNN
model = Sequential()
```

```
# First convolution layer and pooling
model.add(Convolution2D(32, (3, 3), input_shape=(sz, sz, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
# Second convolution layer and pooling
model.add(Convolution2D(32, (3, 3), activation='relu'))
# input_shape is going to be the pooled feature maps from the previous convolu
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
# Flattening the layers
model.add(Flatten())
```

```
# Adding a fully connected layer
model.add(Dense(units=96, activation='relu'))
model.add(Dropout(0.40))
model.add(Dense(units=32, activation='relu'))
model.add(Dense(units=7, activation='softmax')) # softmax for more than 2
```

```
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_2 (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d_2 (MaxPoolin g2D)	(None, 63, 63, 32)	0
conv2d_3 (Conv2D)	(None, 61, 61, 32)	9248
max_pooling2d_3 (MaxPoolin g2D)	(None, 30, 30, 32)	0
flatten_1 (Flatten)	(None, 28800)	0
dense_3 (Dense)	(None, 96)	2764896

dropout_1 (Dropout)	(None, 96)	0
dense_4 (Dense)	(None, 32)	3104
dense_5 (Dense)	(None, 7)	231

```

=====
Total params: 2778375 (10.60 MB)
Trainable params: 2778375 (10.60 MB)
Non-trainable params: 0 (0.00 Byte)

```

```
model.compile(optimizer='adam',loss=tf.keras.losses.SparseCategoricalCrossentropy
```

```

history = model.fit(
    train_generator,
    validation_data=test_generator,
    epochs=35
)

```

```

Epoch 7/35
105/105 [=====] - 116s 1s/step
Epoch 8/35
105/105 [=====] - 114s 1s/step
Epoch 9/35
105/105 [=====] - 115s 1s/step
Epoch 10/35
105/105 [=====] - 155s 1s/step
Epoch 11/35
105/105 [=====] - 114s 1s/step

```

Epoch 18/35
105/105 [=====] - 116s 1s/steps
Epoch 19/35
105/105 [=====] - 114s 1s/steps
Epoch 20/35
105/105 [=====] - 115s 1s/steps
Epoch 21/35
105/105 [=====] - 115s 1s/steps
Epoch 22/35
105/105 [=====] - 114s 1s/steps
Epoch 23/35
105/105 [=====] - 155s 1s/steps
Epoch 24/35
105/105 [=====] - 114s 1s/steps
Epoch 25/35
105/105 [=====] - 116s 1s/steps
Epoch 26/35
105/105 [=====] - 114s 1s/steps
Epoch 27/35
105/105 [=====] - 117s 1s/steps
Epoch 28/35
105/105 [=====] - 114s 1s/steps
Epoch 29/35
105/105 [=====] - 114s 1s/steps
Epoch 30/35
105/105 [=====] - 155s 1s/steps
Epoch 31/35
105/105 [=====] - 114s 1s/steps
Epoch 32/35
105/105 [=====] - 116s 1s/steps
Epoch 33/35
105/105 [=====] - 113s 1s/steps
Epoch 34/35
105/105 [=====] - 114s 1s/steps
Epoch 35/35
105/105 [=====] - 113s 1s/steps

```
scores = model.evaluate(test_generator)
```

105/105 [=====] - 40s 384ms/steps

◀  ▶

scores


```
history.history.keys()
```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
type(history.history['loss'])
```

```
list
```

```
len(history.history['loss'])
```

```
35
```

```
history.history['loss'][:1] # show loss for first 5 epochs
```

```
[1.8676438331604004]
```

```
acc = history.history['accuracy']  
val_acc = history.history['val_accuracy']
```

```
loss = history.history['loss']  
val_loss = history.history['val_loss']
```

```
# import matplotlib.pyplot as plt
# EPOCHS = 20

# plt.figure(figsize=(8, 8))
# plt.subplot(1, 2, 1)
# plt.plot(range(EPOCHS), acc, label='Training Accuracy')
# plt.plot(range(EPOCHS), val_acc, label='Validation Accuracy')
# plt.legend(loc='lower right')
# plt.title('Training and Validation Accuracy')

# plt.subplot(1, 2, 2)
# plt.plot(range(EPOCHS), loss, label='Training Loss')
# plt.plot(range(EPOCHS), val_loss, label='Validation Loss')
# plt.legend(loc='upper right')
# plt.title('Training and Validation Loss')
# plt.show()
```

```
def predict(model, img):
    img_array = tf.keras.preprocessing.image.img_to_array(images[i])
    img_array = tf.expand_dims(img_array, 0)

    predictions = model.predict(img_array)

    predicted_class = class_names[np.argmax(predictions[0])]
    confidence = round(100 * (np.max(predictions[0])), 2)
    return predicted_class, confidence
```

ChatGPT

```
plt.figure(figsize=(15, 15))
for images, labels in test_generator:
    for i in range(6):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i])

        predicted_class, confidence = predict(model, images[i])
        actual_class = class_names[int(labels[i])]

        plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class}.\n Confidence: {confidence}")

        plt.axis("off")
    break
```



1/1 [=====] - 0s 344ms/step
1/1 [=====] - 0s 69ms/step
1/1 [=====] - 0s 54ms/step
1/1 [=====] - 0s 78ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 35ms/step

Actual: Hyundai Creta,
Predicted: Hyundai Creta.
Confidence: 93.93%



Actual: Toyota Innova,
Predicted: Toyota Innova.
Confidence: 99.96%



Actual: Toyota Innova,
Predicted: Toyota Innova.
Confidence: 99.28%



Actual: Toyota Innova,
Predicted: Toyota Innova.
Confidence: 96.71%



Actual: Mahindra Scorpio,
Predicted: Mahindra Scorpio.
Confidence: 94.68%



Actual: Rolls Royce,
Predicted: Rolls Royce.
Confidence: 76.7%

